

Pet Translator

Bridging the Gap Between Pets and Humans

Team: Amine KHALIL • Mohamed MELLOUK • Anas ISARTI • Abir ISLAM

Machine Learning (ML01)

Architecture

An End-to-End Hybrid Pipeline

- **Frontend:** Real-time audio capture via React.
- **Edge Inference:** Local, ultra-fast classification on the smartphone via TensorFlow.js.
- **Cloud Backend:** FastAPI server connecting to a semantic database.
- **Generative AI: Llama 3.2** generating contextual human sentences.

Audio Signal Processing

Mohamed MELLOUK

- **Datasets:** ESC-50, AudioSet, ShivaRao (Dogs), CatMeows (Cats).
- **The Challenge:** Dog sounds are *acoustic types* (bark). Cat sounds are *situational* (food).
- **Noise Reduction:** Applying `noisereduce` to clean real-world ambient sounds.
- **Isolation:** Using Voice Activity Detection (WebRTCvad) to isolate vocalization.

Feature Extraction

Visualizing Sound

- Audio waves are converted into visual representations (**Mel-spectrograms**).
- **Standardization:** Fixed durations of 4 seconds for dogs and 2 seconds for cats.
- **Perception:** Logarithmic scaling applied on Mel frequencies to mimic biological hearing.
- **Result:** 96x96 feature matrices, formatted for Deep Learning.

Classification Models

Anas ISARTI

- **The Problem:** Small datasets (113 dog clips, 440 cat clips) lead to severe overfitting.
- **The Solution:** Transfer Learning. We used a frozen **MobileNetV2** architecture.
- **Deployment:** The Keras model was converted to `.tflite`, quantized to INT8, and deployed seamlessly in the browser.

The Data Ceiling

Analyzing Limits

- The Dog model performs well (**~85% Accuracy**).
- The Cat model plateaued at **~35% F1-score** for the `food` class.
- **Experimentation:** Hyperparameter tuning, Architectures swaps, Class-imbalance handling.
- **Conclusion:** The bottleneck is the dataset size itself, not the model capacity.

Linguistic Brain

Abir ISLAM

- **Objective:** Give a real voice and distinct personality to the pet's raw intent.
- **Dataset:** Created a synthetic dataset mapping [Intent + Context + Personality] to [Sentence] .
- **Fine-Tuning:** Trained **Llama 3.2** using LoRA on Google Colab GPUs.
- **Deployment:** Converted weights to GGUF format for lightning-fast CPU inference.

Grounding the AI

The RAG System

- **The Problem:** LLMs can hallucinate biologically inaccurate translations.
- **The Solution:** Retrieval-Augmented Generation (RAG).
- **Implementation:**
 - Built a veterinary behavioral knowledge base.
 - Embedded using `sentence-transformers` and stored in **ChromaDB**.
 - The LLM retrieves real scientific context before generating the translation.

Frontend Integration

Amine KHALIL

- **User Interface:** A modern, responsive React application.
- **Pet Selector:** Allows the user to select the animal and its specific personality.
- **Real-Time Feedback:** Displays local TFJS probabilities instantly before querying the remote LLM.

Demonstration

(Live Video Demonstration)

Conclusion

- **Deployed** a functional hybrid Edge/Cloud AI system.
- **Applied** rigorous scientific methodology to identify data limitations.
- **Mastered** the full stack: Audio Processing, Computer Vision, and NLP.

Thank you.